

# CS 5220: Homework 2

Group 19: Robert Carson (rac428), Robert Chiodi (rmc298), Sam Tung (sat83)

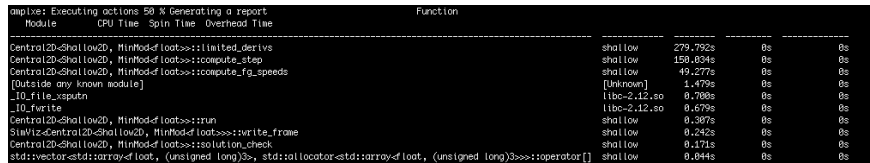
October 14, 2015

## 1 Initial Profiling

In order to perform initial profiling of the code before any improvements are made, Intel's VTUNE was used via the terminal command line on Totient. In an attempt to get the most accurate results (taken with a large sample size), we decided to gather data while running the large wave simulation, invoked with the command `make big`. Information collected from VTUNE's 'advanced-hotspots' option is used in our analysis.

### 1.1 Whole Program - Advanced Hotspots

First, hotspots in the entire program were examined in order to determine where our efforts should be directed. The time taken in the top 10 most time consuming functions can be seen below in Figure 1. Of these, it is clear that most of our optimization efforts should be directed to the functions `limited_derivs`, `compute_step`, and `compute_fg_speeds`. These functions were then examined individually, once again using VTUNE on our advanced-hotspots collection.



Module	CPU Time	Spin Time	Overhead Time	Function
Central2D-Shallow2D, MinModOfloats::limited_derivs	shallow	279.792s	0s	0s
Central2D-Shallow2D, MinModOfloats::compute_step	shallow	150.084s	0s	0s
Central2D-Shallow2D, MinModOfloats::compute_fg_speeds	shallow	49.277s	0s	0s
[Outside any known module]	[Unknown]	1.479s	0s	0s
_IO_file_xsputn	libc-2.12.so	0.700s	0s	0s
_IO_fwrite	libc-2.12.so	0.679s	0s	0s
Central2D-Shallow2D, MinModOfloats::run	shallow	0.597s	0s	0s
SimViz-Central2D-Shallow2D, MinModOfloats::write_frame	shallow	0.242s	0s	0s
Central2D-Shallow2D, MinModOfloats::solution_check	shallow	0.171s	0s	0s
std::vector<std::array<float, (unsigned long)3>, std::allocator<std::array<float, (unsigned long)3>>>::operator[]	shallow	0.044s	0s	0s

Figure 1: Top 10 most time consuming functions in the wave simulation. Generated using Intel's VTUNE on Totient.

### 1.2 `limited_derivs` - Advanced Hotspots

The function `limited_deriv` is used to calculate the fluxes into and out of each cell in order to advance to the next time step. This involves a three point computational stencil in each direction and loops through the entire domain interior (the whole domain except for those where boundary conditions are applied). Each point requires `du.size() × 9` floating point operations as well as `du.size() × 2` calls to the intrinsic function `min`. Sadly, the hotspot

analysis on the `limited_deriv` function, shown in Figure 2, does not give any hints on possible optimizations or bottle necks.

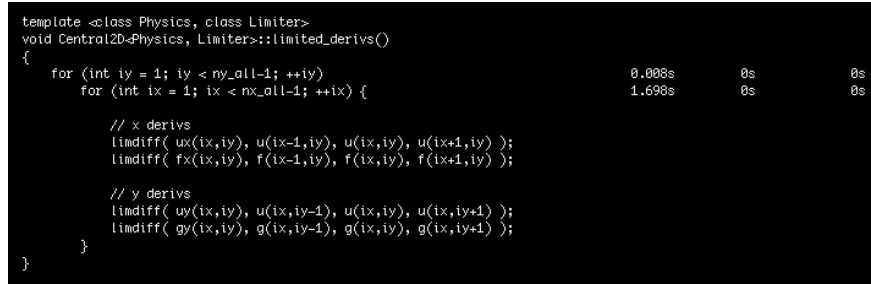


Figure 2: Time taken to perform each loop present in `limited_derivs`

### 1.3 `compute_step` - Advanced Hotspots

The purpose of `compute_step` is to update the wave equation to the next time step using a predictor-corrector method. First, the fluxes are calculated in the prediction. Next, the corrector step uses the predicted fluxes, the differences in velocities, and the current velocities to advance to the next time state. Luckily, VTUNE’s report is more helpful than in the previous case, and provides extensive timings for this function, shown in Figure 3. The calculation in the corrector step can be seen to be the most expensive cost of the function. It is important to note, however, that the predictor step and copying of the solution to the `u` array sum to half of the function’s cost.

### 1.4 `compute_fg_speeds` - Advanced Hotspots

The function of `compute_fg_speeds` has two primary responsibilities: to update the cell centered fluxes, `f` and `g`, and to calculate the maximum speed in the domain, allowing dynamic adjustment of the time step in order to satisfy the CFL condition and ensure numerical stability. The timing data for this function can be seen in Figure 4. While the most time consuming portion of the code is most likely the calculation of the fluxes and wave speed (both of which are in the `Shallow2d` structure), the calls to the intrinsic function `max` also represent a non-trivial amount of time.

template <class Physics, class Limiter>			
void Central2D<Physics, Limiter>::compute_step(int io, real dt)			
{			
real dtcdx2 = 0.5 * dt / dx;			
real dtcdy2 = 0.5 * dt / dy;			
// Predictor (flux values of f and g at half step)			
for (int iy = 1; iy < ny_all-1; ++iy)	0.003s	0s	0s
for (int ix = 1; ix < nx_all-1; ++ix) {	0.357s	0s	0s
vec uh = u(ix,iy);	1.909s	0s	0s
for (int m = 0; m < uh.size(); ++m) {			
uh[m] -= dtcdx2 * fx(ix,iy)[m];			
uh[m] -= dtcdy2 * gy(ix,iy)[m];	7.334s	0s	0s
}			
Physics::flux(f(ix,iy), g(ix,iy), uh);			
}			
// Corrector (finish the step)			
for (int iy = nghost-io; iy < ny+nghost-io; ++iy)	0.010s	0s	0s
for (int ix = nghost-io; ix < nx+nghost-io; ++ix) {	1.750s	0s	0s
for (int m = 0; m < v(ix,iy).size(); ++m) {			
v(ix,iy)[m] =	26.781s	0s	0s
0.2500 * ( u(ix, iy)[m] + u(ix+1,iy ) [m] +			
u(ix,iy+1)[m] + u(ix+1,iy+1)[m] ) -	1.553s	0s	0s
0.0625 * ( ux(ix+1,iy ) [m] - ux(ix,iy ) [m] +	6.718s	0s	0s
ux(ix+1,iy+1)[m] - ux(ix,iy+1)[m] +			
uy(ix, iy+1)[m] - uy(ix, iy) [m] +			
uy(ix+1,iy+1)[m] - uy(ix+1,iy) [m] ) -	1.501s	0s	0s
dtcdx2 * ( f(ix+1,iy ) [m] - f(ix,iy ) [m] +	0.228s	0s	0s
f(ix+1,iy+1)[m] - f(ix,iy+1)[m] ) -	7.972s	0s	0s
dtcdy2 * ( g(ix, iy+1)[m] - g(ix, iy) [m] +	1.840s	0s	0s
g(ix+1,iy+1)[m] - g(ix+1,iy) [m] );	8.207s	0s	0s
}			
}			
// Copy from v storage back to main grid			
for (int j = nghost; j < ny+nghost; ++j){	0.003s	0s	0s
for (int i = nghost; i < nx+nghost; ++i){	1.624s	0s	0s
u(i,j) = v(i-io,j-io);	10.566s	0s	0s
}			
}			
}	0.001s	0s	0s

Figure 3: Time taken to perform each loop present in `limited_derivs`

{			
using namespace std;			
real cx = 1.0e-15;			
real cy = 1.0e-15;			
for (int iy = 0; iy < ny_all; ++iy)	0.001s	0s	0s
for (int ix = 0; ix < nx_all; ++ix) {	3.240s	0s	0s
real cell_cx, cell_cy;			
Physics::flux(f(ix,iy), g(ix,iy), u(ix,iy));			
Physics::wave_speed(cell_cx, cell_cy, u(ix,iy));			
cx = max(cx, cell_cx);	4.299s	0s	0s
cy = max(cy, cell_cy);	1.666s	0s	0s
}			
cx_ = cx;			
cy_ = cy;			
}			

Figure 4: Time taken to perform each loop present in `limited_derivs`

## References